

JIAMU (MORRIS) ZHANG

Houston, TX | 216-280-7699 | mz81@rice.edu | MorrisZJ.github.io | [Google Scholar](#) | [LinkedIn](#) | [GitHub](#)

RESEARCH

Efficient and scalable LLM and agentic systems — *system–method co-design* — and rigorous ways to measure whether the efficiency is real.

EDUCATION

Ph.D. in Computer Science

08/2024 – Present

Rice University, Houston, TX

- Advisor: Prof. Hanjie Chen, Chili Lab (08/2025 – Present)
- Advisor: Prof. Xia “Ben” Hu, Data Lab (08/2024 – 08/2025)

B.S. in Computer Science

05/2024

Case Western Reserve University, Cleveland, OH

- Department of Computer and Data Science. Advisors: Prof. Shuai Xu and Prof. Vipin Chaudhary.

RESEARCH INTERESTS

- **Memory-bound LLM inference.** KV cache compression, quantization and transform coding for long-context and agentic workloads.
- **Conditional computation.** Mixture-of-Experts and training-free routing that turn static redundancy in pretrained models into input-dependent sparsity — and the serving systems that realize it under real memory budgets.
- **Efficient and verifiable agentic systems.** Cutting redundant computation and communication in LLM agents and multi-agent systems — and rigorously evaluating whether reported efficiency gains are real.

EXPERIENCE

Research Intern, AI Research & Development

05/2026 – Present

Nokia, Sunnyvale, CA

- Efficiency and system–method co-design for large-model inference; led the design and open-source release of AnyJev (below).

OPEN SOURCE

AnyJev — lead author and principal maintainer

506★ · 73 forks in 3 days

2026

Nokia Applied Research · github.com/nokia-applied-research/AnyJev · `pip install anyjev`

- Turns any LLM into a typed decision model with calibrated probabilities, training-free by default — served through vLLM as a pooling endpoint plus a few kilobytes of closed-form head.
- Cuts option-order sensitivity 0.230→0.073 and calibration error 0.240→0.095 with zero labels; raises the share auto-decidable at 5% risk from 7.7% to 52.0% with 100–500 labels.

PUBLICATIONS

EMNLP’26

Rethinking the Evaluation of Efficiency Methods for Multi-Agent Systems

Award Nominee

Jiamu Zhang*, Lingxi Zhang*, Pengjun Lu*, Qiyue Zhang, Yu-Neng Chuang, Zhengchen Li, Shuai Xu, Vipin Chaudhary, Hanjie Chen.

Conference on Empirical Methods in Natural Language Processing (EMNLP), Main Conference, 2026. [arXiv]

CVPR’25

Flexible Group Count Enables Hassle-Free Structured Pruning

Jiamu Zhang*, Shaochen Zhong*, Andrew Ye, Zirui Liu, Sebastian Zhao, Kaixiong Zhou, Li Li, Soo-Hyun Choi, Rui Chen, Xia Hu, Shuai Xu, Vipin Chaudhary.

IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2025. Acceptance rate 22.1%. [PDF]

Preprint

SPECTRA: Pushing the KV Cache Beyond the 2-Bit Cliff via Spectral Transform Coding

Jiamu Zhang, Liang Wu, Kelly Wan, Hanjie Chen, Liangjie Hong.

arXiv:2608.07915, 2026. [arXiv]

Preprint	WiSP: A Working-Set View of Mixture-of-Experts Serving on Extremely Low-Resource Hardware Jiamu Zhang , Liang Wu, Mayank Darbari, Liangjie Hong. <i>arXiv:2606.21868</i> , 2026. [arXiv]
In Review	Recognizing and Restructuring Latent Experts for Large Language Model Compression Jiamu Zhang , Alessandro Mason, Ning Xie, Aarav Swami, Ashley Chen, Shuai Xu, Vipin Chaudhary, Hanjie Chen. <i>Under review</i> , 2026.
ACL'25	ReasonerRank: Redefining Language Model Evaluation with Ground-Truth-Free Ranking Frameworks Jiamu Zhang , Jiayi Yuan, Andrew Wen, Hoang Anh Duy Le, Yu-Neng Chuang, Soo-Hyun Choi, Rui Chen, Xia Hu. <i>Findings of the Association for Computational Linguistics: ACL 2025</i> , pp. 13623–13639. [PDF]
NeurIPS'23	One Less Reason for Filter Pruning: Gaining Free Adversarial Robustness with Structured Grouped Kernel Pruning Shaochen Zhong, Jiamu Zhang* , Zaichuan You*, Sebastian Zhao*, Zachary LeClaire, Zirui Liu, Vipin Chaudhary, Shuai Xu, Xia Hu. <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , 2023. [PDF]
arXiv	The Science of Evaluating Foundation Models Jiayi Yuan, Jiamu Zhang , Andrew Wen, Xia Hu. <i>arXiv:2502.09670</i> , 2025. [arXiv]
TMLR'25	Stop Overthinking: A Survey on Efficient Reasoning for Large Language Models 589 citations · 787★ Yang Sui, Yu-Neng Chuang, Guanchu Wang, Jiamu Zhang , Tianyi Zhang, Jiayi Yuan, Hongyi Liu, Andrew Wen, Shaochen Zhong, Hanjie Chen, Xia Hu. <i>Transactions on Machine Learning Research (TMLR)</i> , 2025. [arXiv] [Code]
ICML'24	GNNs Also Deserve Editing, and They Need It More Than Once Shaochen Zhong*, Hoang Anh Duy Le*, Zirui Liu, Zhimeng Jiang, Andrew Ye, Jiamu Zhang , Jiayi Yuan, Kaixiong Zhou, Zhaozhuo Xu, Jing Ma, Shuai Xu, Vipin Chaudhary, Xia Hu. <i>International Conference on Machine Learning (ICML)</i> , 2024. [PDF]
npj Med	Context Matching is not Reasoning: Assessing Generalized Evaluation of Generative Language Models in Clinical Settings Andrew Wen, Qiuha Lu, Yu-Neng Chuang, Guanchu Wang, Jiayi Yuan, Jiamu Zhang , Liwei Wang, Sunyang Fu, Kurt D. Miller, Heling Jia, Steven D. Bedrick, William R. Hersh, Kirk E. Roberts, Xia Hu, Hongfang Liu. <i>npj Digital Medicine</i> , 2025.

* denotes equal contribution. Citation, star and download counts as of September 2026.

HONORS AND AWARDS

EMNLP 2026 Award Nomination — for “Rethinking the Evaluation of Efficiency Methods for Multi-Agent Systems”	2026
Swanger Graduate Fellowship — Case Western Reserve University	2024
The Computer and Data Science Research Award — Case Western Reserve University	2024
Dean’s High Honor List & Latin Honors — Case Western Reserve University	2024

PROFESSIONAL SERVICE

Conference Reviewer — NeurIPS, ICML, ACL
Journal Reviewer — IEEE TPAMI, IEEE TCDS, ACM TIST, ACM TCH, Neurocomputing
Languages — English, Mandarin Chinese

TEACHING

TA, COMP 584 — Natural Language Processing · Rice University	2026 Spring
TA, COMP 631 — Introduction to Information Retrieval · Rice University	2025 Spring
TA, CSDS 302 — Discrete Mathematics · Case Western Reserve University	2022, 2023, 2024
TA, CSDS 386 — Quantum Computing, Information, and Devices · Case Western Reserve University	2023 Spring