

# JIAMU (MORRIS) ZHANG

Houston, TX | 216-280-7699 | mz81@rice.edu | MorrisZJ.github.io | Google Scholar | LinkedIn | GitHub

## RESEARCH

Efficient and scalable LLM and agentic systems — *system–method co-design* — and rigorous ways to measure whether the efficiency is real.

## EDUCATION

### Ph.D. in Computer Science

08/2024 – Present

Rice University, Houston, TX

- Advisor: Prof. Hanjie Chen, Chili Lab (08/2025 – Present)
- Advisor: Prof. Xia “Ben” Hu, Data Lab (08/2024 – 08/2025)

### B.S. in Computer Science

05/2024

Case Western Reserve University, Cleveland, OH

- Department of Computer and Data Science. Advisors: Prof. Shuai Xu and Prof. Vipin Chaudhary.

## RESEARCH INTERESTS

- **Memory-bound LLM inference.** KV cache compression, quantization and transform coding for long-context and agentic workloads.
- **Conditional computation.** Mixture-of-Experts and training-free routing that turn static redundancy in pretrained models into input-dependent sparsity — and the serving systems that realize it under real memory budgets.
- **Efficient and verifiable agentic systems.** Cutting redundant computation and communication in LLM agents and multi-agent systems — and rigorously evaluating whether reported efficiency gains are real.

## EXPERIENCE

### Research Intern, AI Research & Development

05/2026 – Present

Nokia, Sunnyvale, CA

- Efficiency and system–method co-design for large-model inference; led the design and open-source release of AnyJev (below).

## OPEN SOURCE

### AnyJev — lead author and principal maintainer

726★ · 96 forks

2026

Nokia Applied Research · [github.com/nokia-applied-research/AnyJev](https://github.com/nokia-applied-research/AnyJev) · `pip install anyjev`

- Turns any LLM into a typed decision model with calibrated probabilities, training-free by default — served through vLLM as a pooling endpoint plus a few kilobytes of closed-form head.
- Cuts option-order sensitivity 0.230→0.073 and calibration error 0.240→0.095 with zero labels; raises the share auto-decidable at 5% risk from 7.7% to 52.0% with 100–500 labels.
- Passed 500 GitHub stars in its first three days; released under Apache-2.0 with a published PyPI package.

## PUBLICATIONS

EMNLP’26

### Rethinking the Evaluation of Efficiency Methods for Multi-Agent Systems

Award Nominee

Jiamu Zhang\*, Lingxi Zhang\*, Pengjun Lu\*, Qiyue Zhang, Yu-Neng Chuang, Zhengchen Li, Shuai Xu, Vipin Chaudhary, Hanjie Chen.

Conference on Empirical Methods in Natural Language Processing (EMNLP), Main Conference, Oral, 2026. [arXiv]

CVPR’25

### Flexible Group Count Enables Hassle-Free Structured Pruning

Jiamu Zhang\*, Shaochen Zhong\*, Andrew Ye, Zirui Liu, Sebastian Zhao, Kaixiong Zhou, Li Li, Soo-Hyun Choi, Rui Chen, Xia Hu, Shuai Xu, Vipin Chaudhary.

IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2025. Acceptance rate 22.1%. [PDF]

Preprint

### SPECTRA: Pushing the KV Cache Beyond the 2-Bit Cliff via Spectral Transform Coding

Jiamu Zhang, Liang Wu, Kelly Wan, Hanjie Chen, Liangjie Hong.

arXiv:2608.07915, 2026. [arXiv]

|            |                                                                                                                                                                                                                                                                                                                                                                                                |
|------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Preprint   | <b>WiSP: A Working-Set View of Mixture-of-Experts Serving on Extremely Low-Resource Hardware</b><br><b>Jiamu Zhang</b> , Liang Wu, Mayank Darbari, Liangjie Hong.<br><i>arXiv:2606.21868</i> , 2026. <a href="#">[arXiv]</a>                                                                                                                                                                   |
| In Review  | <b>Recognizing and Restructuring Latent Experts for Large Language Model Compression</b><br><b>Jiamu Zhang</b> , Alessandro Mason, Ning Xie, Aarav Swami, Ashley Chen, Shuai Xu, Vipin Chaudhary, Hanjie Chen.<br><i>Under review</i> , 2026.                                                                                                                                                  |
| ACL'25     | <b>ReasonerRank: Redefining Language Model Evaluation with Ground-Truth-Free Ranking Frameworks</b><br><b>Jiamu Zhang</b> , Jiayi Yuan, Andrew Wen, Hoang Anh Duy Le, Yu-Neng Chuang, Soo-Hyun Choi, Rui Chen, Xia Hu.<br><i>Findings of the Association for Computational Linguistics: ACL 2025</i> , pp. 13623–13639. <a href="#">[PDF]</a>                                                  |
| NeurIPS'23 | <b>One Less Reason for Filter Pruning: Gaining Free Adversarial Robustness with Structured Grouped Kernel Pruning</b><br>Shaochen Zhong, <b>Jiamu Zhang*</b> , Zaichuan You*, Sebastian Zhao*, Zachary LeClaire, Zirui Liu, Vipin Chaudhary, Shuai Xu, Xia Hu.<br><i>Advances in Neural Information Processing Systems (NeurIPS)</i> , 2023. <a href="#">[PDF]</a>                             |
| arXiv      | <b>The Science of Evaluating Foundation Models</b><br>Jiayi Yuan, <b>Jiamu Zhang</b> , Andrew Wen, Xia Hu.<br><i>arXiv:2502.09670</i> , 2025. <a href="#">[arXiv]</a>                                                                                                                                                                                                                          |
| TMLR'25    | <b>Stop Overthinking: A Survey on Efficient Reasoning for Large Language Models</b> <span>589 citations · 787★</span><br>Yang Sui, Yu-Neng Chuang, Guanchu Wang, <b>Jiamu Zhang</b> , Tianyi Zhang, Jiayi Yuan, Hongyi Liu, Andrew Wen, Shaochen Zhong, Hanjie Chen, Xia Hu.<br><i>Transactions on Machine Learning Research (TMLR)</i> , 2025. <a href="#">[arXiv]</a> <a href="#">[Code]</a> |
| ICML'24    | <b>GNNs Also Deserve Editing, and They Need It More Than Once</b><br>Shaochen Zhong*, Hoang Anh Duy Le*, Zirui Liu, Zhimeng Jiang, Andrew Ye, <b>Jiamu Zhang</b> , Jiayi Yuan, Kaixiong Zhou, Zhaozhuo Xu, Jing Ma, Shuai Xu, Vipin Chaudhary, Xia Hu.<br><i>International Conference on Machine Learning (ICML)</i> , 2024. <a href="#">[PDF]</a>                                             |
| npj Med    | <b>Context Matching is not Reasoning: Assessing Generalized Evaluation of Generative Language Models in Clinical Settings</b><br>Andrew Wen, Qiuha Lu, Yu-Neng Chuang, Guanchu Wang, Jiayi Yuan, <b>Jiamu Zhang</b> , Liwei Wang, Sunyang Fu, Kurt D. Miller, Heling Jia, Steven D. Bedrick, William R. Hersh, Kirk E. Roberts, Xia Hu, Hongfang Liu.<br><i>npj Digital Medicine</i> , 2025.   |

\* denotes equal contribution. Citation, star and download counts as of September 2026.

## HONORS AND AWARDS

|                                                                                                             |      |
|-------------------------------------------------------------------------------------------------------------|------|
| EMNLP 2026 Award Nomination — for “Rethinking the Evaluation of Efficiency Methods for Multi-Agent Systems” | 2026 |
| Swanger Graduate Fellowship — Case Western Reserve University                                               | 2024 |
| The Computer and Data Science Research Award — Case Western Reserve University                              | 2024 |
| Dean’s High Honor List & Latin Honors — Case Western Reserve University                                     | 2024 |

## PROFESSIONAL SERVICE

|                                                                             |
|-----------------------------------------------------------------------------|
| Conference Reviewer — NeurIPS, ICML, ACL                                    |
| Journal Reviewer — IEEE TPAMI, IEEE TCDS, ACM TIST, ACM TCH, Neurocomputing |
| Languages — English, Mandarin Chinese                                       |

## TEACHING

|                                                                                              |                  |
|----------------------------------------------------------------------------------------------|------------------|
| TA, COMP 584 — Natural Language Processing · Rice University                                 | 2026 Spring      |
| TA, COMP 631 — Introduction to Information Retrieval · Rice University                       | 2025 Spring      |
| TA, CSDS 302 — Discrete Mathematics · Case Western Reserve University                        | 2022, 2023, 2024 |
| TA, CSDS 386 — Quantum Computing, Information, and Devices · Case Western Reserve University | 2023 Spring      |