JIAMU (MORRIS) ZHANG

Houston, TX | 216-280-7699 | mz81@rice.edu | MorrisZJ.github.io | Google Scholar | LinkedIn | GitHub

RESEARCH

Efficient and scalable LLM and agentic systems — *system–method co-design* — and rigorous ways to measure whether the efficiency is real.

EDUCATION

Ph.D. in Computer Science 08/2024 - Present

Rice University, Houston, TX

- Advisor: Prof. Hanjie Chen, Chili Lab (08/2025 – Present)
- Advisor: Prof. Xia "Ben" Hu, Data Lab (08/2024 – 08/2025)

B.S. in Computer Science 05/2024

Case Western Reserve University, Cleveland, OH

- Department of Computer and Data Science. Advisors: Prof. Shuai Xu and Prof. Vipin Chaudhary.

RESEARCH INTERESTS

- **Memory-bound LLM inference.** KV cache compression, quantization and transform coding for long-context and agentic workloads.
- **Conditional computation.** Mixture-of-Experts and training-free routing that turn static redundancy in pretrained models into input-dependent sparsity — and the serving systems that realize it under real memory budgets.
- **Efficient and verifiable agentic systems.** Cutting redundant computation and communication in LLM agents and multi-agent systems — and rigorously evaluating whether reported efficiency gains are real.

EXPERIENCE

Research Intern, AI Research & Development 05/2026 - Present

Nokia, Sunnyvale, CA

- Efficiency and system–method co-design for large-model inference; led the design and open-source release of AnyJev (below).

OPEN SOURCE

AnyJev — lead author and principal maintainer 1,003★ · 131 forks 2026

Nokia Applied Research · github.com/nokia-applied-research/AnyJev · `pip install anyjev` · technical report arXiv:2610.00831

- Turns any LLM into a typed decision model with calibrated probabilities, training-free by default — served through vLLM as a pooling endpoint plus a few kilobytes of closed-form head.
- Cuts option-order sensitivity 0.230→0.073 and calibration error 0.240→0.095 with zero labels; raises the share auto-decidable at 5% risk from 7.7% to 52.0% with 100–500 labels.
- Passed 1,000 GitHub stars in its first ten days; released under Apache-2.0 with a published PyPI package.
- Joint 4th of 148 systems on the community Jev Decision Index (v0.3.2), the only training-free entry in the top band.

PUBLICATIONS

EMNLP'26 **Rethinking the Evaluation of Efficiency Methods for Multi-Agent Systems** Award Nominee
Jiamu Zhang*, Lingxi Zhang*, Pengjun Lu*, Qiyue Zhang, Yu-Neng Chuang, Zhengchen Li, Shuai Xu, Vipin Chaudhary, Hanjie Chen.
Conference on Empirical Methods in Natural Language Processing (EMNLP), Main Conference, Oral, 2026. [arXiv]

CVPR'25 **Flexible Group Count Enables Hassle-Free Structured Pruning**
Jiamu Zhang*, Shaochen Zhong*, Andrew Ye, Zirui Liu, Sebastian Zhao, Kaixiong Zhou, Li Li, Soo-Hyun Choi, Rui Chen, Xia Hu, Shuai Xu, Vipin Chaudhary.
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2025. Acceptance rate 22.1%. [PDF]

Preprint **SPECTRA: Pushing the KV Cache Beyond the 2-Bit Cliff via Spectral Transform Coding**
Jiamu Zhang, Liang Wu, Kelly Wan, Hanjie Chen, Liangjie Hong.
arXiv:2608.07915, 2026. [arXiv]

Preprint **WiSP: A Working-Set View of Mixture-of-Experts Serving on Extremely Low-Resource Hardware**
Jiamu Zhang, Liang Wu, Mayank Darbari, Liangjie Hong.
arXiv:2606.21868, 2026. [arXiv]

In Review **BellowsKV: One Latent Cache for Many Layers, Expanded on Demand**
Jiamu Zhang, Hanjie Chen.
Under review, 2026.

In Review **Recognizing and Restructuring Latent Experts for Large Language Model Compression**
Jiamu Zhang, Alessandro Mason, Ning Xie, Aarav Swami, Ashley Chen, Shuai Xu, Vipin Chaudhary, Hanjie Chen.
Under review, 2026.

ACL'25 **ReasonerRank: Redefining Language Model Evaluation with Ground-Truth-Free Ranking Frameworks**
Jiamu Zhang, Jiayi Yuan, Andrew Wen, Hoang Anh Duy Le, Yu-Neng Chuang, Soo-Hyun Choi, Rui Chen, Xia Hu.
Findings of the Association for Computational Linguistics: ACL 2025, pp. 13623–13639. [PDF]

NeurIPS'23 **One Less Reason for Filter Pruning: Gaining Free Adversarial Robustness with Structured Grouped Kernel Pruning**
Shaochen Zhong, **Jiamu Zhang***, Zaichuan You*, Sebastian Zhao*, Zachary LeClaire, Zirui Liu, Vipin Chaudhary, Shuai Xu, Xia Hu.
Advances in Neural Information Processing Systems (NeurIPS), 2023. [PDF]

arXiv **The Science of Evaluating Foundation Models**
Jiayi Yuan, **Jiamu Zhang**, Andrew Wen, Xia Hu.
arXiv:2502.09670, 2025. [arXiv]

TMLR'25 **Stop Overthinking: A Survey on Efficient Reasoning for Large Language Models** 603 citations · 791★
Yang Sui, Yu-Neng Chuang, Guanchu Wang, **Jiamu Zhang**, Tianyi Zhang, Jiayi Yuan, Hongyi Liu, Andrew Wen, Shaochen Zhong, Hanjie Chen, Xia Hu.
Transactions on Machine Learning Research (TMLR), 2025. [arXiv] [Code]

ICML'24 **GNNs Also Deserve Editing, and They Need It More Than Once**
Shaochen Zhong*, Hoang Anh Duy Le*, Zirui Liu, Zhimeng Jiang, Andrew Ye, **Jiamu Zhang**, Jiayi Yuan, Kaixiong Zhou, Zhaozhuo Xu, Jing Ma, Shuai Xu, Vipin Chaudhary, Xia Hu.
International Conference on Machine Learning (ICML), 2024. [PDF]

npj Med **Context Matching is not Reasoning: Assessing Generalized Evaluation of Generative Language Models in Clinical Settings**
Andrew Wen, Qiuhao Lu, Yu-Neng Chuang, Guanchu Wang, Jiayi Yuan, **Jiamu Zhang**, Liwei Wang, Sunyang Fu, Kurt D. Miller, Heling Jia, Steven D. Bedrick, William R. Hersh, Kirk E. Roberts, Xia Hu, Hongfang Liu.
npj Digital Medicine, 2025.

** denotes equal contribution. Citation counts from Google Scholar; star and download counts from GitHub and PyPI. Figures as of October 2026.*

HONORS AND AWARDS

EMNLP 2026 Award Nomination — for "Rethinking the Evaluation of Efficiency Methods for Multi-Agent Systems" 2026

Swanger Graduate Fellowship — Case Western Reserve University 2024

The Computer and Data Science Research Award — Case Western Reserve University 2024

Dean's High Honor List & Latin Honors — Case Western Reserve University 2024

PROFESSIONAL SERVICE

Conference Reviewer — NeurIPS, ICML, ACL

Journal Reviewer — IEEE TPAMI, IEEE TCDS, ACM TIST, ACM TCH, Neurocomputing

Languages — English, Mandarin Chinese

TEACHING

Guest Lecture, COMP 640 — Graduate Seminar in ML: Large Language Models · Rice University 2026 Fall
"A Simple Intro to Test-Time Scaling & Reasoning: Thinking longer, thinking better"

TA, COMP 584 — Natural Language Processing · Rice University 2026 Spring

TA, COMP 631 — Introduction to Information Retrieval · Rice University 2025 Spring

TA, CSDS 302 — Discrete Mathematics · Case Western Reserve University 2022, 2023, 2024

TA, CSDS 386 — Quantum Computing, Information, and Devices · Case Western Reserve University 2023 Spring